

With PyTorch 0.2.0_2


PYTORCH

Lab I I

CNN MNIST

Sung Kim <hunkim+ml@gmail.com>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



With PyTorch 0.2.0_2




Lab 11-1

CNN Basics

Sung Kim <hunkim+ml@gmail.com>


Code: <https://github.com/hunkim/DeepLearningZeroToAll/>




<https://github.com/hunkim/DeepLearningZeroToAll/tree/master/pytorch>

https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-11-1-mnist_cnn.py

CNN




CNN for CT images



Asan Medical Center & Microsoft Medical Bigdata Contest Winner by GeunYoung Lee and Alex Kim

<https://www.slideshare.net/GYLee3/ss-72966495>

Convolution layer and max pooling



Single depth slice


1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters
and stride 2


6	8
3	4

Simple convolution layer

Stride: 1x1



Simple CNN



Conv layer I

```
class CNN(torch.nn.Module):
```

```
    def __init__(self):
```

```
        super(CNN, self).__init__()
```

```
        # L1 ImgIn shape=(?, 28, 28, 1)
```

```
        #   Conv    -> (?, 28, 28, 32)
```

```
        #   Pool    -> (?, 14, 14, 32)
```

```
        self.layer1 = torch.nn.Sequential(
```

```
            torch.nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),
```

```
            torch.nn.ReLU(),
```

```
            torch.nn.MaxPool2d(kernel_size=2, stride=2))
```

```
        # L2 ImgIn shape=(?, 14, 14, 32)
```

```
        #   Conv    -> (?, 14, 14, 64)
```

```
        #   Pool    -> (?, 7, 7, 64)
```

```
        self.layer2 = torch.nn.Sequential(
```

```
            torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
```


```
            torch.nn.ReLU(),
```

```
            torch.nn.MaxPool2d(kernel_size=2, stride=2))
```

```
        # Final FC 7x7x64 inputs -> 10 outputs
```


```
        self.fc = torch.nn.Linear(7 * 7 * 64, 10, bias=True)
```

```
        torch.nn.init.xavier_uniform(self.fc.weight)
```



https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-11-1-mnist_cnn.py

Conv layer 2



```
class CNN(torch.nn.Module):
```

```
    def __init__(self):
```

```
        super(CNN, self).__init__()
```

```
        # L1 ImgIn shape=(?, 28, 28, 1)
```

```
        #   Conv    -> (?, 28, 28, 32)
```

```
        #   Pool    -> (?, 14, 14, 32)
```

```
        self.layer1 = torch.nn.Sequential(
```

```
            torch.nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),
```

```
            torch.nn.ReLU(),
```

```
            torch.nn.MaxPool2d(kernel_size=2, stride=2))
```

```
        # L2 ImgIn shape=(?, 14, 14, 32)
```

```
        #   Conv    -> (?, 14, 14, 64)
```

```
        #   Pool    -> (?, 7, 7, 64)
```

```
        self.layer2 = torch.nn.Sequential(
```

```
            torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
```

```
            torch.nn.ReLU(),
```

```
            torch.nn.MaxPool2d(kernel_size=2, stride=2))
```

```
        # Final FC 7x7x64 inputs -> 10 outputs
```


```
        self.fc = torch.nn.Linear(7 * 7 * 64, 10, bias=True)
```

```
        torch.nn.init.xavier_uniform(self.fc.weight)
```

https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-11-1-mnist_cnn.py

Fully Connected (FC, Dense) layer

```
class CNN(torch.nn.Module):  
  
    def __init__(self):  
        super(CNN, self).__init__()  
        # L1 ImgIn shape=(?, 28, 28, 1)  
        #  Conv    -> (?, 28, 28, 32)  
        #  Pool    -> (?, 14, 14, 32)  
        self.layer1 = torch.nn.Sequential(  
            torch.nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),  
            torch.nn.ReLU(),  
            torch.nn.MaxPool2d(kernel_size=2, stride=2))  
        # L2 ImgIn shape=(?, 14, 14, 32)  
        #  Conv    -> (?, 14, 14, 64)  
        #  Pool    -> (?, 7, 7, 64)  
        self.layer2 = torch.nn.Sequential(  
            torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),  
            torch.nn.ReLU(),  
            torch.nn.MaxPool2d(kernel_size=2, stride=2))  
  
        # Final FC 7x7x64 inputs -> 10 outputs  
        self.fc = torch.nn.Linear(7 * 7 * 64, 10, bias=True)  
        torch.nn.init.xavier_uniform(self.fc.weight)
```



Training and Evaluation

```
# instantiate CNN model
model = CNN()

# define cost/loss & optimizer
criterion = torch.nn.CrossEntropyLoss()    # Softmax is internally computed.
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

# train my model
print('Learning started. It takes sometime.')
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = len(mnist_train) // batch_size

    for i, (batch_xs, batch_ys) in enumerate(data_loader):
        X = Variable(batch_xs)    # image is already size of (28x28), no reshape
        Y = Variable(batch_ys)    # label is not one-hot encoded

        optimizer.zero_grad()
        hypothesis = model(X)
        cost = criterion(hypothesis, Y)
        cost.backward()
        optimizer.step()

        avg_cost += cost.data / total_batch

    print("[Epoch: {:>4}] cost = {:>.9}".format(epoch + 1, avg_cost[0]))

print('Learning Finished!')
```

```

# instantiate CNN model
model = CNN()

# define cost/loss & optimizer
criterion = torch.nn.CrossEntropyLoss() # Softmax is internally computed.
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

# train my model
print('Learning started. It takes sometime.')
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = len(mnist_train) // batch_size

    for i, (batch_xs, batch_ys) in enumerate(data_loader):
        X = Variable(batch_xs) # image is already size of (28x28), no reshape
        Y = Variable(batch_ys) # label is not one-hot encoded

        optimizer.zero_grad()
        hypothesis = model(X)
        cost = criterion(hypothesis, Y)
        cost.backward()
        optimizer.step()

    avg_cost += cost.data / total_batch

    print("[Epoch: {:>4}] cost = {:>.9}".format(epoch + 1, avg_cost[0]))

print('Learning Finished!')

```

Training and Evaluation

Epoch: 0001 cost = 0.340291267
 Epoch: 0002 cost = 0.090731326
 Epoch: 0003 cost = 0.064477619
 Epoch: 0004 cost = 0.050683064
 ...
 Epoch: 0011 cost = 0.017758641
 Epoch: 0012 cost = 0.014156652
 Epoch: 0013 cost = 0.012397016
 Epoch: 0014 cost = 0.010693789
 Epoch: 0015 cost = 0.009469977
 Learning Finished!
 Accuracy: **0.9885**

With PyTorch 0.2.0_2


PYTORCH

Lab 11-2

Deep CNN

Sung Kim <hunkim+ml@gmail.com>


Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/tree/master/pytorch>

https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-11-2-mnist_deep_cnn.py

Deep CNN



Deep CNN

```
class CNN(torch.nn.Module):

    def __init__(self):
        super(CNN, self).__init__()
        # L1 ImgIn shape=(?, 28, 28, 1)
        #   Conv    -> (?, 28, 28, 32)
        #   Pool    -> (?, 14, 14, 32)
        self.layer1 = torch.nn.Sequential(
            torch.nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2),
            torch.nn.Dropout(p=1 - keep_prob))
        # L2 ImgIn shape=(?, 14, 14, 32)
        #   Conv    -> (?, 14, 14, 64)
        #   Pool    -> (?, 7, 7, 64)
        self.layer2 = torch.nn.Sequential(
            torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2),
            torch.nn.Dropout(p=1 - keep_prob))
        # L3 ImgIn shape=(?, 7, 7, 64)
        #   Conv    -> (?, 7, 7, 128)
        #   Pool    -> (?, 4, 4, 128)
        self.layer3 = torch.nn.Sequential(
            torch.nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2, padding=1),
            torch.nn.Dropout(p=1 - keep_prob))
```

```
# L4 FC 4x4x128 inputs -> 625 outputs
self.fc1 = torch.nn.Linear(4 * 4 * 128, 625, bias=True)
torch.nn.init.xavier_uniform(self.fc1.weight)
self.layer4 = torch.nn.Sequential(
    self.fc1,
    torch.nn.ReLU(),
    torch.nn.Dropout(p=1 - keep_prob))
# L5 Final FC 625 inputs -> 10 outputs
self.fc2 = torch.nn.Linear(625, 10, bias=True)
torch.nn.init.xavier_uniform(self.fc2.weight)
```

```
def forward(self, x):
    out = self.layer1(x)
    out = self.layer2(out)
    out = self.layer3(out)
    out = out.view(out.size(0), -1) # Flatten them for FC
    out = self.fc1(out)
    out = self.fc2(out)
    return out
```

Deep CNN

```
# instantiate CNN model
model = CNN()

# define cost/loss & optimizer
criterion = torch.nn.CrossEntropyLoss() # Softmax is internally computed.
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

# train my model
print('Learning started. It takes sometime.')
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = len(mnist_train) // batch_size

    for i, (batch_xs, batch_ys) in enumerate(data_loader):
        X = Variable(batch_xs) # image is already size of (28x28), no reshape
        Y = Variable(batch_ys) # label is not one-hot encoded

        optimizer.zero_grad()
        hypothesis = model(X)
        cost = criterion(hypothesis, Y)
        cost.backward()
        optimizer.step()

    avg_cost += cost.data / total_batch

    print("[Epoch: {:>4}] cost = {:>.9}".format(epoch + 1, avg_cost[0]))

print('Learning Finished!')
```

```
# Test model and check accuracy
model.eval() # set the model to evaluation mode (dropout=False)

X_test = Variable(mnist_test.test_data.view(len(mnist_test), 1, 28, 28).float())
Y_test = Variable(mnist_test.test_labels)

prediction = model(X_test)
correct_prediction = (torch.max(prediction.data, 1)[1] == Y_test.data)
accuracy = correct_prediction.float().mean()
print('Accuracy:', accuracy)
```

Epoch: 0013 cost = 0.027188021
Epoch: 0014 cost = 0.023604777
Epoch: 0015 cost = 0.024607201
Learning Finished!

Accuracy: 0.9938

With PyTorch 0.2.0_2


PYTORCH

Lab 11-3

CNN Class

Sung Kim <hunkim+ml@gmail.com>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/tree/master/pytorch>

https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-11-3-mnist_cnn_class.py

```
class CNN(torch.nn.Module):
```

```
def __init__(self):
    super(CNN, self).__init__()
    self._build_net()

def _build_net(self):
    # dropout (keep_prob) rate 0.7~0.5 on training, but should be 1
    self.keep_prob = 0.7

    # L1 ImgIn shape=(?, 28, 28, 1)
    #  Conv -> (?, 28, 28, 32)
    #  Pool -> (?, 14, 14, 32)
    self.layer1 = torch.nn.Sequential(
        torch.nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),
        torch.nn.ReLU(),
        torch.nn.MaxPool2d(kernel_size=2, stride=2),
        torch.nn.Dropout(p=1 - self.keep_prob))

    # L2 ImgIn shape=(?, 14, 14, 32)
    #  Conv -> (?, 14, 14, 64)
    #  Pool -> (?, 7, 7, 64)
    self.layer2 = torch.nn.Sequential(
        torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
        torch.nn.ReLU(),
        torch.nn.MaxPool2d(kernel_size=2, stride=2),
        torch.nn.Dropout(p=1 - self.keep_prob))

    # L3 ImgIn shape=(?, 7, 7, 64)
    #  Conv -> (?, 7, 7, 128)
    #  Pool -> (?, 4, 4, 128)
    self.layer3 = torch.nn.Sequential(
        torch.nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
        torch.nn.ReLU(),
        torch.nn.MaxPool2d(kernel_size=2, stride=2, padding=1),
        torch.nn.Dropout(p=1 - self.keep_prob))
```

CNN

Python Class

```
# L4 FC 4x4x128 inputs -> 625 outputs
self.keep_prob = 0.5
self.fc1 = torch.nn.Linear(4 * 4 * 128, 625, bias=True)
torch.nn.init.xavier_uniform(self.fc1.weight)
self.layer4 = torch.nn.Sequential(
    self.fc1,
    torch.nn.ReLU(),
    torch.nn.Dropout(p=1 - self.keep_prob))

# L5 Final FC 625 inputs -> 10 outputs
self.fc2 = torch.nn.Linear(625, 10, bias=True)
torch.nn.init.xavier_uniform(self.fc2.weight)

# define cost/loss & optimizer
self.criterion = torch.nn.CrossEntropyLoss() # Softmax is internally computed.
self.optimizer = torch.optim.Adam(self.parameters(), lr=learning_rate)
```

CNN

Python Class

```
def forward(self, x):
    out = self.layer1(x)
    out = self.layer2(out)
    out = self.layer3(out)
    out = out.view(out.size(0), -1)  # Flatten them for FC
    out = self.fc1(out)
    out = self.fc2(out)
    return out

def predict(self, x):
    self.eval()
    return self.forward(x)

def get_accuracy(self, x, y):
    prediction = self.predict(x)
    correct_prediction = (torch.max(prediction.data, 1)[1] == y.data)
    self.accuracy = correct_prediction.float().mean()
    return self.accuracy

def train_model(self, x, y):
    self.train()
    self.optimizer.zero_grad()
    hypothesis = self.forward(x)
    self.cost = self.criterion(hypothesis, y)
    self.cost.backward()
    self.optimizer.step()
    return self.cost
```

```
# instantiate CNN model
model = CNN()

# train my model
print('Learning started. It takes sometime.')
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = len(mnist_train) // batch_size

    for i, (batch_xs, batch_ys) in enumerate(data_loader):
        X = Variable(batch_xs)  # image is already size of (28x28), no reshape
        Y = Variable(batch_ys)  # label is not one-hot encoded

        cost = model.train_model(X, Y)

        avg_cost += cost.data / total_batch

    print("[Epoch: {:>4}] cost = {:.>.9}".format(epoch + 1, avg_cost[0]))

print('Learning Finished!')

# Test model and check accuracy
X_test = Variable(mnist_test.test_data.view(len(mnist_test), 1, 28, 28).float())
Y_test = Variable(mnist_test.test_labels)

print('Accuracy:', model.get_accuracy(X_test, Y_test))
```


Exercise

- Deep & Wide?
- CIFAR 10
- ImageNet

Here are the classes in the dataset, as well as 10 random images from each


airplane




automobile



bird



cat




deer




dog



frog



horse



ship




truck



With PyTorch 0.2.0_2

PYTORCH



Lab 12

RNN

Sung Kim <hunkim+ml@gmail.com>
<http://hunkim.github.io/ml/>