



With PyTorch 0.2.0\_2

**PYTORCH**




Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



With PyTorch 0.2.0\_2

**PYTORCH**

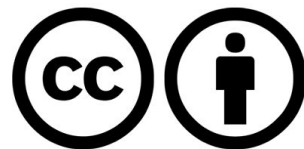


# Lab 12-1

Hi Hello RNN

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>


Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/tree/master/pytorch>

<https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-12-1-hello-rnn.py>

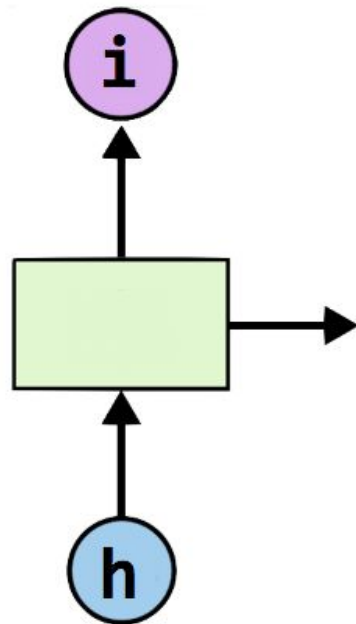
# Teach RNN 'hihello'



- text: 'hihello'
- unique chars (vocabulary, voc):  
h, i, e, l, o
- voc index:  
h:0, i:1, e:2, l:3, o:4


# One-hot encoding

[1, 0, 0, 0, 0],	# h 0
[0, 1, 0, 0, 0],	# i 1
[0, 0, 1, 0, 0],	# e 2
[0, 0, 0, 1, 0],	# l 3
[0, 0, 0, 0, 1],	# o 4




# Teach RNN 'hihello'

$[1, 0, 0, 0, 0]$ ,	# h 0
$[0, 1, 0, 0, 0]$ ,	# i 1
$[0, 0, 1, 0, 0]$ ,	# e 2
$[0, 0, 0, 1, 0]$ ,	# l 3
$[0, 0, 0, 0, 1]$ ,	# o 4



# Teach RNN 'hihello'

[1, 0, 0, 0, 0],	# h 0
[0, 1, 0, 0, 0],	# i 1
[0, 0, 1, 0, 0],	# e 2
[0, 0, 0, 1, 0],	# l 3
[0, 0, 0, 0, 1],	# o 4

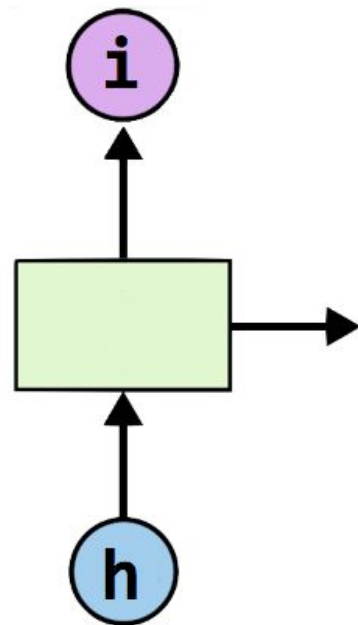


# Teach RNN 'hihello'

```
idx2char = ['h', 'i', 'e', 'l', 'o']

# Teach hello: hihell -> ihello
x_data = [[0, 1, 0, 2, 3, 3]] # hihell
x_one_hot = [[[1, 0, 0, 0, 0], # h 0
               [0, 1, 0, 0, 0], # i 1
               [1, 0, 0, 0, 0], # h 0
               [0, 0, 1, 0, 0], # e 2
               [0, 0, 0, 1, 0], # l 3
               [0, 0, 0, 1, 0]]] # l 3

y_data = [1, 0, 2, 3, 3, 4] # ihello
```






# RNN parameters

```
# As we have one batch of samples, we will change them to variables only once
inputs = torch.Tensor(x_one_hot)
labels = torch.LongTensor(y_data)

inputs = Variable(inputs)
labels = Variable(labels)

num_classes = 5
input_size = 5 # one-hot size
hidden_size = 5 # output from the LSTM. 5 to directly predict one-hot
batch_size = 1 # one sentence
sequence_length = 6 # |ihello| == 6
num_layers = 1 # one-layer rnn
```



# Class RNN

```
class RNN(nn.Module):
```

```
    def __init__(self, num_classes, input_size, hidden_size, num_layers):
        super(RNN, self).__init__()
        self.num_classes = num_classes
        self.num_layers = num_layers
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.sequence_length = sequence_length
        # Set parameters for RNN block
        # Note: batch_first=False by default.
        # When true, inputs are (batch_size, sequence_length, input_dimension)
        # instead of (sequence_length, batch_size, input_dimension)
        self.rnn = nn.RNN(input_size=input_size, hidden_size=hidden_size,
                           num_layers=num_layers, batch_first=True)
        # Fully connected layer to obtain outputs corresponding to the number
        # of classes
        self.fc = nn.Linear(hidden_size, num_classes)
```

```
    def forward(self, x):
```

```
        # Initialize hidden and cell states
        h_0 = Variable(torch.zeros(
            x.size(0), self.num_layers, self.hidden_size))

        # Reshape input
        x = x.view(x.size(0), self.sequence_length, self.input_size)

        # Propagate input through RNN
        # Input: (batch, seq_len, input_size)
        # h_0: (batch, num_layers * num_directions, hidden_size)

        out, _ = self.rnn(x, h_0)

        # Reshape output from (batch, seq_len, hidden_size) to (batch *
        # seq_len, hidden_size)
        out = out.view(-1, self.hidden_size)
        # Return outputs applied to fully connected layer
        out = self.fc(out)
        return out
```

# Loss, Optimizer

```
# Instantiate RNN model
rnn = RNN(num_classes, input_size, hidden_size, num_layers)


# Set loss and optimizer function
criterion = torch.nn.CrossEntropyLoss()    # Softmax is internally computed.
optimizer = torch.optim.Adam(rnn.parameters(), lr=learning_rate)
```

# Training

```
# Train the model
for epoch in range(num_epochs):
    outputs = rnn(inputs)
    optimizer.zero_grad()
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()
    _, idx = outputs.max(1)
    idx = idx.data.numpy()
    result_str = [idx2char[c] for c in idx.squeeze()]
    print("epoch: %d, loss: %1.3f" % (epoch + 1, loss.data[0]))
    print("Predicted string: ", ''.join(result_str))
```

With PyTorch 0.2.0\_2

**PYTORCH**




# Lab 12-2

RNN with long sequences

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/tree/master/pytorch>

<https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-12-2-char-seq-rnn.py>

# Manual data creation

```
idx2char = ['h', 'i', 'e', 'l', 'o']
x_data = [[0, 1, 0, 2, 3, 3]]      # hihell
x_one_hot = [[[1, 0, 0, 0, 0],      # h 0
               [0, 1, 0, 0, 0],      # i 1
               [1, 0, 0, 0, 0],      # h 0
               [0, 0, 1, 0, 0],      # e 2
               [0, 0, 0, 1, 0],      # l 3
               [0, 0, 0, 1, 0]]]     # l 3
y_data = [[1, 0, 2, 3, 3, 4]]     # ihello
```

# Better data creation

```
sample = " if you want you"  
idx2char = list(set(sample)) # index -> char  
char2idx = {c: i for i, c in enumerate(idx2char)} # char -> idx
```



# Hyper parameters

```
sample = " if you want you"
idx2char = list(set(sample)) # index -> char
char2idx = {c: i for i, c in enumerate(idx2char)} # char -> idx

# hyper parameters
dic_size = len(char2idx) # RNN input size (one hot size)
rnn_hidden_size = len(char2idx) # RNN output size
num_classes = len(char2idx) # final output size (RNN or softmax, etc.)
batch_size = 1 # one sample data, one batch
sequence_length = len(sample) - 1 # number of lstm unfolding (unit #)
```

# One hot encoding

```
sample_idx = [char2idx[c] for c in sample] # char to index
x_data = [sample_idx[:-1]] # X data sample (0 ~ n-1) hello: hell
y_data = [sample_idx[1:]] # Y label sample (1 ~ n) hello: ello
```

```
x_data = torch.Tensor(x_data)
y_data = torch.LongTensor(y_data)
```

```
# one hot encoding
```

```
def one_hot(x, num_classes):
    idx = x.long()
    idx = idx.view(-1, 1)
    x_one_hot = torch.zeros(x.size()[0] * x.size()[1], num_classes)
    x_one_hot.scatter_(1, idx, 1)
    x_one_hot = x_one_hot.view(x.size()[0], x.size()[1], num_classes)
    return x_one_hot
```

```
x_one_hot = one_hot(x_data, num_classes)
```

```
inputs = Variable(x_one_hot)
labels = Variable(y_data)
```

<https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-12-2-char-seq-rnn.py>

# LSTM

```
class LSTM(nn.Module):
```

```
    def __init__(self, num_classes, input_size, hidden_size, num_layers):
        super(LSTM, self).__init__()
        self.num_classes = num_classes
        self.num_layers = num_layers
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.sequence_length = sequence_length
        # Set parameters for RNN block
        # Note: batch_first=False by default.
        # When true, inputs are (batch_size, sequence_length, input_dimension)
        # instead of (sequence_length, batch_size, input_dimension)
        self.lstm = nn.LSTM(input_size=input_size, hidden_size=hidden_size,
                            num_layers=num_layers, batch_first=True)
        # Fully connected layer to obtain outputs corresponding to the number
        # of classes
        self.fc = nn.Linear(hidden_size, num_classes)
```

```
    def forward(self, x):
```

```
        # Initialize hidden and cell states
        h_0 = Variable(torch.zeros(
            self.num_layers, x.size(0), self.hidden_size))
        c_0 = Variable(torch.zeros(
            self.num_layers, x.size(0), self.hidden_size))

        # Reshape input
        x.view(x.size(0), self.sequence_length, self.input_size)
```

```
        # Propagate input through RNN
        # Input: (batch, seq_len, input_size)
        # h_0: (num_layers * num_directions, batch, hidden_size)
        out, _ = self.lstm(x, (h_0, c_0))
```

```
        # Reshape output from (batch, seq_len, hidden_size) to (batch *
        # seq_len, hidden_size)
        out = out.view(-1, self.hidden_size)
        # Return outputs applied to fully connected layer
        out = self.fc(out)
        return out
```

# Loss, Optimizer

```
# Instantiate RNN model
rnn = RNN(num_classes, input_size, hidden_size, num_layers)

# Set loss and optimizer function
criterion = torch.nn.CrossEntropyLoss()    # Softmax is internally computed.
optimizer = torch.optim.Adam(rnn.parameters(), lr=learning_rate)
```


# Training

```
# Train the model
for epoch in range(num_epochs):
    outputs = lstm(inputs)
    optimizer.zero_grad()
    loss = criterion(outputs, labels.view(-1))
    loss.backward()
    optimizer.step()
    _, idx = outputs.max(1)
    idx = idx.data.numpy()
    result_str = [idx2char[c] for c in idx.squeeze()]
    print("epoch: %d, loss: %1.3f" % (epoch + 1, loss.data[0]))
    print("Predicted string: ", ''.join(result_str))

print("Learning finished!")
```


With PyTorch 0.2.0\_2

**PYTORCH**



Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>

Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/tree/master/pytorch>

[https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-12-4-rnn-long\\_char.py](https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-12-4-rnn-long_char.py)

# Really long sentence?

```
sentence = ("if you want to build a ship, don't drum up people together to "  
            "collect wood and don't assign them tasks and work, but rather "  
            "teach them to long for the endless immensity of the sea.")
```



# Really long sentence?

```
sentence = ("if you want to build a ship, don't drum up people together to "  
            "collect wood and don't assign them tasks and work, but rather "  
            "teach them to long for the endless immensity of the sea.")
```

*# training dataset*

0 if you wan -> f you want

1 f you want -> you want

2 you want -> you want t

3 you want t -> ou want to

...

168 of the se -> of the sea

169 of the sea -> f the sea.

# RNN parameters

```
char_set = list(set(sentence))
char_dic = {w: i for i, w in enumerate(char_set)}

# hyperparameters
learning_rate = 0.1
num_epochs = 500
input_size = len(char_set) # RNN input size (one hot size)
hidden_size = len(char_set) # RNN output size
num_classes = len(char_set) # final output size (RNN or softmax, etc.)
sequence_length = 10 # any arbitrary number
num_layers = 2 # number of layers in RNN
```

## *# training dataset*

0 if you wan -> f you want  
1 f you want -> you want  
2 you want -> you want t  
3 you want t -> ou want to  
...  
168 of the se -> of the sea  
169 of the sea -> f the sea.

```
char_set = list(set(sentence))
char_dic = {w: i for i, w in enumerate(char_set)}
```

```
dataX = []
dataY = []
```

```
for i in range(0, len(sentence) - seq_length):
    x_str = sentence[i:i + seq_length]
    y_str = sentence[i + 1: i + seq_length + 1]
    print(i, x_str, '->', y_str)
```

```
x = [char_dic[c] for c in x_str] # x str to index
y = [char_dic[c] for c in y_str] # y str to index
```

```
dataX.append(x)
dataY.append(y)
```

# Making dataset

*# training dataset*

0 if you wan -> f you want

1 f you want -> you want

2 you want -> you want t


3 you want t -> ou want to

...

168 of the se -> of the sea

169 of the sea -> f the sea.

# Wide & Deep



# Stacked RNN


```
# one hot encoding
```

```
def one_hot(x, num_classes):  
    idx = x.long()  
    idx = idx.view(-1, 1)  
    x_one_hot = torch.zeros(x.size()[0] * x.size()[1], num_classes)  
    x_one_hot.scatter_(1, idx, 1)  
    x_one_hot = x_one_hot.view(x.size()[0], x.size()[1], num_classes)  
    return x_one_hot
```

```
x_one_hot = one_hot(x_data, num_classes)
```

```
inputs = Variable(x_one_hot)
```

```
labels = Variable(y_data)
```



# LSTM

```
class LSTM(nn.Module):
```


```
    def __init__(self, num_classes, input_size, hidden_size, num_layers):
        super(LSTM, self).__init__()
        self.num_classes = num_classes
        self.num_layers = num_layers
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.sequence_length = sequence_length
        # Set parameters for RNN block
        # Note: batch_first=False by default.
        # When true, inputs are (batch_size, sequence_length, input_dimen:
        # instead of (sequence_length, batch_size, input_dimension)
        self.lstm = nn.LSTM(input_size=input_size, hidden_size=hidden_size,
                            num_layers=num_layers, batch_first=True)
        # Fully connected layer
        self.fc = nn.Linear(hidden_size, num_classes)
```

```
    def forward(self, x):
```


```
        # Initialize hidden and cell states
        h_0 = Variable(torch.zeros(
            self.num_layers, x.size(0), self.hidden_size))
        c_0 = Variable(torch.zeros(
            self.num_layers, x.size(0), self.hidden_size))
        # h_0 = Variable(torch.zeros(
        # self.num_layers, x.size(0), self.hidden_size))
        # c_0 = Variable(torch.zeros(
        # self.num_layers, x.size(0), self.hidden_size))

        # Propagate input through LSTM
        # Input: (batch, seq_len, input_size)
        out, _ = self.lstm(x, (h_0, c_0))
        # Note: the output tensor of LSTM in this case is a block with holes
        # > add .contiguous() to apply view()
        out = out.contiguous().view(-1, self.hidden_size)
        # Return outputs applied to fully connected layer
        out = self.fc(out)
        return out
```

# Softmax (FC) in Deep CNN



# Softmax






# Loss, Optimizer

```
# Instantiate RNN model
lstm = LSTM(num_classes, input_size, hidden_size, num_layers)

# Set loss and optimizer function
criterion = torch.nn.CrossEntropyLoss()    # Softmax is internally computed.
optimizer = torch.optim.Adam(lstm.parameters(), lr=learning_rate)
```



# Training


```
# Train the model
for epoch in range(num_epochs):
    outputs = lstm(inputs)
    optimizer.zero_grad()
    # obtain the loss function
    # flatten target labels to match output
    loss = criterion(outputs, labels.view(-1))
    loss.backward()
    optimizer.step()
    # obtain the predicted indices of the next character
    _, idx = outputs.max(1)
    idx = idx.data.numpy()
    idx = idx.reshape(-1, sequence_length) # (170,10)
    # display the prediction of the last sequence
    result_str = [char_set[c] for c in idx[-1]]
    print("epoch: %d, loss: %1.3f" % (epoch + 1, loss.data[0]))
    print("Predicted string: ", ''.join(result_str))

print("Learning finished!")
```

# Print results

**g you want to build a ship, don't drum up people together to collect wood and don't assign them tasks and work, but rather teach them to long for the endless immensity of the sea.**

# char-rnn



## Shakespeare

It looks like we can learn to spell English words. But how about if there is more structure and style in the data? To examine this I downloaded all the works of Shakespeare and concatenated them into a single (4.4MB) file. We can now afford to train a larger network, in this case lets try a 3-layer RNN with 512 hidden nodes on each layer. After we train the network for a few hours we obtain samples such as:

PANDARUS:

Alas, I think he shall be come approached and the day  
When little strain would be attain'd into being never fed,  
And who is but a chain and subjects of his death,  
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,  
Breaking and strongly should be buried, when I perish  
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and  
my fair nudes begun out of the fact, to be conveyed,  
Whose noble souls I'll have the heart of the wars.

Clown:

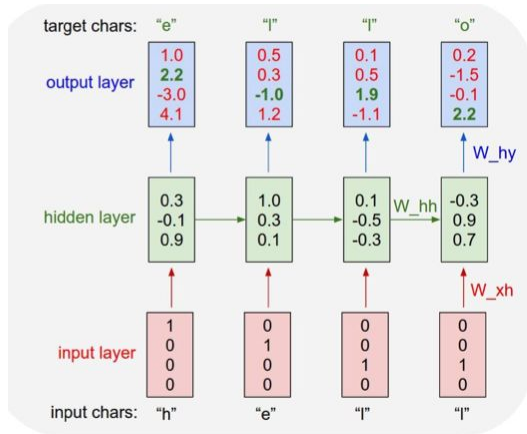
Come, sir, I will make did behold your worship.

VIOLA:

I'll drink it.

## Linux Source Code

I wanted to push structured data to its limit, so for the final challenge I decided to use code. In particular, I took all the source and header files found in the [Linux repo on Github](#), concatenated all of them in a single giant file (474MB of C code) (I was originally going to train only on the kernel but that by itself is only ~16MB). Then I trained several as-large-as-fits-on-my-GPU 3-layer LSTMs over a period of a few days. These models have about 10 million parameters, which is still on the lower end for RNN models. The results are superfun:



```
/*
 * Increment the size file of the new incorrect UI_FILTER group information
 * of the size generatively.
 */
static int indicate_policy(void)
{
    int error;
    if (fd == MARN_EPT) {
        /*
         * The kernel blank will coeld it to userspace.
         */
        if (ss->segment < mem_total)
            unblock_graph_and_set_blocked();
        else
            ret = 1;
        goto bail;
    }
    segaddr = in_SB(in.addr);
    selector = seg / 16;
    setup_works = true;
    for (i = 0; i < blocks; i++) {
        seq = buf[i++];
        bpf = bd->bd.next + i * search;
        if (fd) {
            current = blocked;
        }
    }
    rw->name = "Getjbbregs";
    bprm_self_clearl(&iv->version);
    regs->new = blocks[(BPF_STATS << info->historidac)] | PFMR_CLOBATHINC_SECONDS << 12;
    return segtable;
}
```

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

With TF 1.0!

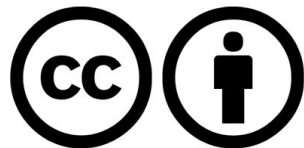


# Lab 12-5

## RNN with time series data (stock)

Sung Kim <[hunkim+ml@gmail.com](mailto:hunkim+ml@gmail.com)>


Code: <https://github.com/hunkim/DeepLearningZeroToAll/>



<https://github.com/hunkim/DeepLearningZeroToAll/tree/master/pytorch>

[https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-12-5-stock\\_prediction.py](https://github.com/hunkim/DeepLearningZeroToAll/blob/master/pytorch/lab-12-5-stock_prediction.py)

# Time series data







# Time series data

Open	High	Low	Volume	Close
828.659973	833.450012	828.349976	1247700	831.659973
823.02002	828.070007	821.655029	1597800	828.070007
819.929993	824.400024	818.97998	1281700	824.159973
819.359985	823	818.469971	1304000	818.97998
819	823	816	1053600	820.450012
816	820.958984	815.48999	1198100	819.23999
811.700012	815.25	809.780029	1129100	813.669983
809.51001	810.659973	804.539978	989700	809.559998
807	811.840027	803.190002	1155300	808.380005

'data-02-stock\_daily.csv'

# Many to one





Open	High	Low	Volume	Close
828.659973	833.450012	828.349976	1247700	831.659973
823.02002	828.070007	821.655029	1597800	828.070007
819.929993	824.400024	818.97998	1281700	824.159973
819.359985	823	818.469971	1304000	818.97998
819	823	816	1053600	820.450012
816	820.958984	815.48999	1198100	819.23999
811.700012	815.25	809.780029	1129100	813.669983
809.51001	810.659973	804.539978	989700	?
807	811.840027	803.190002	1155300	?

```
timesteps = seq_length = 7
```

```
input_size = 5
```

```
num_layers = 1 # number of layers in RNN
```

```
# Open, High, Low, Close, Volume
```

```
xy = np.loadtxt('data-02-stock_daily.csv', delimiter=',')
```

```
xy = xy[::-1] # reverse order (chronically ordered)
```

```
xy = MinMaxScaler(xy)
```

```
x = xy
```

```
y = xy[:, [-1]] # Close as Label
```

```
dataX = []
```

```
dataY = []
```

```
for i in range(0, len(y) - seq_length):
```

```
    _x = x[i:i + seq_length]
```

```
    _y = y[i + seq_length] # Next close price
```

```
    print(_x, "->", _y)
```

```
    dataX.append(_x)
```

```
    dataY.append(_y)
```

# Reading data

```
[ 0.18667876  0.20948057  0.20878184  0.  
0.21744815]
```

```
[ 0.30697388  0.31463414  0.21899367  
0.01247647  0.21698189]
```

```
[ 0.21914211  0.26390721  0.2246864  
0.45632338  0.22496747]
```

```
[ 0.23312993  0.23641916  0.16268272  
0.57017119  0.14744274]
```

```
[ 0.13431201  0.15175877  0.11617252  
0.39380658  0.13289962]
```

```
[ 0.13973232  0.17060429  0.15860382  
0.28173344  0.18171679]
```

```
[ 0.18933069  0.20057799  0.19187983  
0.29783096  0.2086465 ]]
```

```
-> [ 0.14106001]
```

# Train/Test split

*# split to train and testing*

```
train_size = int(len(dataY) * 0.7)
```

```
test_size = len(dataY) - train_size
```

```
trainX = torch.Tensor(np.array(dataX[0:train_size]))
```

```
trainX = Variable(trainX)
```

```
testX = torch.Tensor(np.array(dataX[train_size:len(dataX)]))
```

```
testX = Variable(testX)
```


```
trainY = torch.Tensor(np.array(dataY[0:train_size]))
```

```
trainY = Variable(trainY)
```

```
testY = torch.Tensor(np.array(dataY[train_size:len(dataY)]))
```

```
testY = Variable(testY)
```

# LSTM



```
class LSTM(nn.Module):
```

```
    def __init__(self, num_classes, input_size, hidden_size, num_layers):
        super(LSTM, self).__init__()
        self.num_classes = num_classes
        self.num_layers = num_layers
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.seq_length = seq_length
        # Set parameters for RNN block
        # Note: batch_first=False by default.
        # When true, inputs are (batch_size, sequence_length, input_dimension)
        # instead of (sequence_length, batch_size, input_dimension)
        self.lstm = nn.LSTM(input_size=input_size, hidden_size=hidden_size,
                             num_layers=num_layers, batch_first=True)
        # Fully connected layer
        self.fc = nn.Linear(hidden_size, num_classes)
```


```
    def forward(self, x):
        # Initialize hidden and cell states
        h_0 = Variable(torch.zeros(
            self.num_layers, x.size(0), self.hidden_size))
        c_0 = Variable(torch.zeros(
            self.num_layers, x.size(0), self.hidden_size))

        # Propagate input through LSTM
        _, (h_out, _) = self.lstm(x, (h_0, c_0))
        h_out = h_out.view(-1, self.hidden_size)
        out = self.fc(h_out)
        return out
```

# Loss, Optimizer

```
# Instantiate RNN model
lstm = LSTM(num_classes, input_size, hidden_size, num_layers)

# Set loss and optimizer function
criterion = torch.nn.MSELoss()    # mean-squared error for regression
optimizer = torch.optim.Adam(lstm.parameters(), lr=learning_rate)
```



```

# Train the model
for epoch in range(num_epochs):
    outputs = lstm(trainX)
    optimizer.zero_grad()
    # obtain the loss function
    loss = criterion(outputs, trainY)
    loss.backward()
    optimizer.step()
    print("Epoch: %d, loss: %1.5f" % (epoch, loss.data[0]))


print("Learning finished!")

# Test the model
lstm.eval()
test_predict = lstm(testX)

# Plot predictions
test_predict = test_predict.data.numpy()
testY = testY.data.numpy()
plt.plot(testY)
plt.plot(test_predict)
plt.xlabel("Time Period")
plt.ylabel("Stock Price")
plt.show()

```

# Training and Results





# Exercise

- Implement stock prediction using linear regression only
- Improve results using more features such as keywords and/or sentiments in top news

# Other RNN applications

- Language Modeling
- Speech Recognition
- Machine Translation
- Conversation Modeling/Question Answering
- Image/Video Captioning
- Image/Music/Dance Generation